

In [1]:

```
import pandas as pd
import numpy as np
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn import preprocessing
import tools as t
```

In [2]:

```
%pylab inline
import scipy.stats
from collections import defaultdict # default dictionary
plt.style.use('ggplot')
matplotlib.rcParams['figure.figsize'] = (10.0, 8.0)

import numpy as np
```

Populating the interactive namespace from numpy and matplotlib


In [3]:

```
flights = pd.read_csv("flights.csv")
flights.head()
```

Out[3]:

|   | YEAR | MONTH | DAY | DAY_OF_WEEK | AIRLINE | FLIGHT_NUMBER | TAIL_NUMBER | ORIGIN |
|---|------|-------|-----|-------------|---------|---------------|-------------|--------|
| 0 | 2015 | 1     | 1   | 4           | AS      | 98            | N407AS      |        |
| 1 | 2015 | 1     | 1   | 4           | AA      | 2336          | N3KUAA      |        |
| 2 | 2015 | 1     | 1   | 4           | US      | 840           | N171US      |        |
| 3 | 2015 | 1     | 1   | 4           | AA      | 258           | N3HYAA      |        |
| 4 | 2015 | 1     | 1   | 4           | AS      | 135           | N527AS      |        |

5 rows × 34 columns



In [4]:

```
airlines = pd.read_csv("airlines.csv")
```

In [5]:

```
airlines
```

Out[5]:

|    | IATA_CODE | AIRLINE                      |
|----|-----------|------------------------------|
| 0  | UA        | United Air Lines Inc.        |
| 1  | AA        | American Airlines Inc.       |
| 2  | US        | US Airways Inc.              |
| 3  | F9        | Frontier Airlines Inc.       |
| 4  | B6        | JetBlue Airways              |
| 5  | OO        | Skywest Airlines Inc.        |
| 6  | AS        | Alaska Airlines Inc.         |
| 7  | NK        | Spirit Air Lines             |
| 8  | WN        | Southwest Airlines Co.       |
| 9  | DL        | Delta Air Lines Inc.         |
| 10 | EV        | Atlantic Southeast Airlines  |
| 11 | HA        | Hawaiian Airlines Inc.       |
| 12 | MQ        | American Eagle Airlines Inc. |
| 13 | VX        | Virgin America               |

In [6]:

```
airports = pd.read_csv("airports.csv")
```

In [7]:

```
airports
```

Out[7]:

|     | IATA_CODE | AIRPORT                             | CITY                           | STATE | COUNTRY | LATITUDE | LC |
|-----|-----------|-------------------------------------|--------------------------------|-------|---------|----------|----|
| 0   | ABE       | Lehigh Valley International Airport | Allentown                      | PA    | USA     | 40.65236 |    |
| 1   | ABI       | Abilene Regional Airport            | Abilene                        | TX    | USA     | 32.41132 |    |
| 2   | ABQ       | Albuquerque International Sunport   | Albuquerque                    | NM    | USA     | 35.04022 | -  |
| 3   | ABR       | Aberdeen Regional Airport           | Aberdeen                       | SD    | USA     | 45.44906 |    |
| 4   | ABY       | Southwest Georgia Regional Airport  | Albany                         | GA    | USA     | 31.53552 |    |
| ... | ...       | ...                                 | ...                            | ...   | ...     | ...      |    |
| 317 | WRG       | Wrangell Airport                    | Wrangell                       | AK    | USA     | 56.48433 | -  |
| 318 | WYS       | Westerly State Airport              | West Yellowstone               | MT    | USA     | 44.68840 | .  |
| 319 | XNA       | Northwest Arkansas Regional Airport | Fayetteville/Springdale/Rogers | AR    | USA     | 36.28187 |    |
| 320 | YAK       | Yakutat Airport                     | Yakutat                        | AK    | USA     | 59.50336 | -  |
| 321 | YUM       | Yuma International Airport          | Yuma                           | AZ    | USA     | 32.65658 | -  |

322 rows × 7 columns

In [8]:

```
# flights['DepDate'] = pd.to_datetime(flights.YEAR*10000+flights.MONTH*100+flights.DAY,
format='%Y%m%d')
```


In [9]:

```
flights.head()
```

Out[9]:

|   | YEAR | MONTH | DAY | DAY_OF_WEEK | AIRLINE | FLIGHT_NUMBER | TAIL_NUMBER | ORIGIN |
|---|------|-------|-----|-------------|---------|---------------|-------------|--------|
| 0 | 2015 | 1     | 1   | 4           | AS      | 98            | N407AS      |        |
| 1 | 2015 | 1     | 1   | 4           | AA      | 2336          | N3KUAA      |        |
| 2 | 2015 | 1     | 1   | 4           | US      | 840           | N171US      |        |
| 3 | 2015 | 1     | 1   | 4           | AA      | 258           | N3HYAA      |        |
| 4 | 2015 | 1     | 1   | 4           | AS      | 135           | N527AS      |        |

5 rows × 34 columns



## Cleaning DATA

In [10]:

```
for dataset in flights:
    flights.loc[flights['ARRIVAL_DELAY'] <= -3, 'Status'] = 0
    flights.loc[flights['ARRIVAL_DELAY'] >= -3, 'Status'] = 1
    flights.loc[flights['ARRIVAL_DELAY'] >= 39, 'Status'] = 2
    flights.loc[flights['DIVERTED'] == 1, 'Status'] = 3
    flights.loc[flights['CANCELLED'] == 1, 'Status'] = 4

#Cancellation cause
flights.loc[flights["CANCELLATION_REASON"] == "A", 'CANCELLATION_REASON'] = "0"
flights.loc[flights["CANCELLATION_REASON"] == "B", 'CANCELLATION_REASON'] = "1"
flights.loc[flights["CANCELLATION_REASON"] == "C", 'CANCELLATION_REASON'] = "2"
flights.loc[flights["CANCELLATION_REASON"] == "D", 'CANCELLATION_REASON'] = "3"

# dropping the unwanted data
# flights = flights.drop("Unnamed: 0",1) #Empty
flights = flights.drop("YEAR",1) #Converted to date
flights = flights.drop("DAY",1) #Converted to date
flights = flights.drop("MONTH",1) #Converted to date
flights = flights.drop("DAY_OF_WEEK",1) #Converted to date
flights = flights.drop("DEPARTURE_TIME",1) #Of the departure data we only keep the expected
flights = flights.drop("DEPARTURE_DELAY",1)
flights = flights.drop("ARRIVAL_TIME",1)
flights = flights.drop("SCHEDULED_TIME",1)
flights = flights.drop("ELAPSED_TIME",1)
flights = flights.drop("AIR_TIME",1)
flights = flights.drop("DIVERTED",1)
flights = flights.drop("CANCELLED",1)
flights = flights.drop("DISTANCE",1)
flights = flights.drop("FLIGHT_NUMBER",1)
flights = flights.drop("TAIL_NUMBER",1)
```

## Converting All columns to numaric value

In [11]:

```

for i in airlines.index:
    flights.AIRLINE.loc[flights[flights.AIRLINE == airlines.loc[i].IATA_CODE].index] = i

for i in airports.index:
    flights.ORIGIN_AIRPORT.loc[flights[flights.ORIGIN_AIRPORT == airports.loc[i].IATA_CODE].index] = i
    flights.DESTINATION_AIRPORT.loc[flights[flights.DESTINATION_AIRPORT == airports.loc[i].IATA_CODE].index] = i

```

C:\Users\Waqar\_Shakeel\anaconda3\lib\site-packages\pandas\core\indexing.p

y:670: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
 self.\_setitem\_with\_indexer(indexer, value)

In [12]:

```
flights
```

Out[12]:

|         | AIRLINE | ORIGIN_AIRPORT | DESTINATION_AIRPORT | SCHEDULED_DEPARTURE | TA  |
|---------|---------|----------------|---------------------|---------------------|-----|
| 0       | 6       | 17             | 277                 | 5                   |     |
| 1       | 1       | 176            | 235                 | 10                  |     |
| 2       | 2       | 278            | 66                  | 20                  |     |
| 3       | 1       | 176            | 203                 | 20                  |     |
| 4       | 6       | 277            | 17                  | 25                  |     |
| ...     | ...     | ...            | ...                 | ...                 | ... |
| 1048570 | 6       | 121            | 277                 | 500                 |     |
| 1048571 | 4       | 248            | 166                 | 500                 |     |
| 1048572 | 4       | 285            | 39                  | 500                 |     |
| 1048573 | 4       | 285            | 166                 | 500                 |     |
| 1048574 | 9       | 109            | 216                 | 500                 |     |

1048575 rows × 20 columns

In [13]:

```
flights = flights.replace(np.nan,9)
```

In [14]:

```
Y = flights.Y.values
```

In [15]:

```
f11 = flights.drop('Y', axis=1)
```

In [16]:

```
x = f11.values
```

## Applying PCA for Dimension Reduction

In [17]:

```
pca = PCA()  
principalComponents = pca.fit_transform(x)
```

## Splitting into Test and Train

In [18]:

```
X_train, X_test, y_train, y_test = train_test_split(principalComponents, Y, test_size=  
0.2, random_state=0)
```

## Selecting Best Hyperparameters

In [19]:

```
maxi = -1
index = -1
for j in range(20):
    if(j == 0):
        continue
    clf = RandomForestClassifier(max_depth=j, random_state=0)
    clf.fit(X_train, y_train)
    p = clf.predict(X_test)
    count = 0
    for i in range(len(p)):
        if(p[i] == y_test[i]):
            count = count + 1
    accu = count/len(p)
    if(maxi < accu):
        maxi = accu
        index = j
print(j, count/len(p))
```

```
1 0.9121188279331474
2 0.9490642061845839
3 0.9898528955964047
4 0.9938201845361562
5 0.9985074982714637
6 0.9992847435805736
7 0.9996137615335098
8 0.9997520444412655
9 0.9998474119638557
10 0.9998664854683738
11 0.9998998641012803
12 0.9999284743580573
13 0.9999189376057983
14 0.9999141692296688
15 0.9999284743580573
16 0.9999284743580573
17 0.9999046324774098
18 0.9999141692296688
19 0.9999189376057983
```

In [20]:

```
clf = RandomForestClassifier(max_depth=15, random_state=0)
clf.fit(X_train, y_train)
```

Out[20]:

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=12, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100,
                        n_jobs=None, oob_score=False, random_state=0, verbose=0,
                        warm_start=False)
```

In [21]:

```
p = clf.predict(X_test)
count = 0
for i in range(len(p)):
    if(p[i]==y_test[i]):
        count = count + 1
print(count/len(p))
```

0.9999284743580573

In [22]:

```
t.print_confusion_matrix(p,y_test)
```

Out[22]:

| Predicted | 0     | 1   | 2    | 3      | All    |
|-----------|-------|-----|------|--------|--------|
| Actual    |       |     |      |        |        |
| 0         | 79427 | 0   | 0    | 12     | 79439  |
| 1         | 0     | 494 | 0    | 0      | 494    |
| 2         | 0     | 0   | 7286 | 0      | 7286   |
| 3         | 3     | 0   | 0    | 122493 | 122496 |
| All       | 79430 | 494 | 7286 | 122505 | 209715 |

In [ ]:

## Will flight arrive on time or not?

In [40]:

```
f12 = flights.drop('ON_TIME', axis=1)
```




In [41]:

f12

Out[41]:

|         | AIRLINE | ORIGIN_AIRPORT | DESTINATION_AIRPORT | SCHEDULED_DEPARTURE | TAJ |
|---------|---------|----------------|---------------------|---------------------|-----|
| 0       | 6       | 17             | 277                 | 5                   |     |
| 1       | 1       | 176            | 235                 | 10                  |     |
| 2       | 2       | 278            | 66                  | 20                  |     |
| 3       | 1       | 176            | 203                 | 20                  |     |
| 4       | 6       | 277            | 17                  | 25                  |     |
| ...     | ...     | ...            | ...                 | ...                 | ... |
| 1048570 | 6       | 121            | 277                 | 500                 |     |
| 1048571 | 4       | 248            | 166                 | 500                 |     |
| 1048572 | 4       | 285            | 39                  | 500                 |     |
| 1048573 | 4       | 285            | 166                 | 500                 |     |
| 1048574 | 9       | 109            | 216                 | 500                 |     |

1048575 rows × 19 columns



In [42]:

x = f12.values

In [43]:

Y = flights.ON\_TIME.values

In [44]:

```
pca = PCA()
principalComponents = pca.fit_transform(x)
```

In [45]:

```
X_train, X_test, y_train, y_test = train_test_split(principalComponents, Y, test_size=
0.2, random_state=0)
```

In [46]:

```

maxi = -1
index = -1
for j in range(20):
    if(j == 0):
        continue
    clf = RandomForestClassifier(max_depth=j, random_state=0)
    clf.fit(X_train, y_train)
    p = clf.predict(X_test)
    count = 0
    for i in range(len(p)):
        if(p[i] == y_test[i]):
            count = count + 1
    accu = count/len(p)
    if(maxi < accu):
        maxi = accu
        index = j
    print(j, count/len(p))

```

```

1 0.947848270271559
2 0.9850654459623776
3 0.9905919938964786
4 0.9964523281596452
5 0.9979257563836635
6 0.9986028657940539
7 0.9991416922966884
8 0.9994707102496245
9 0.9996280666618983
10 0.9997234341844885
11 0.9997615811935245
12 0.9998235700832082
13 0.9998474119638557
14 0.9998426435877262
15 0.9998617170922443
16 0.9998855589728918
17 0.9998807905967623
18 0.9999094008535393
19 0.9998903273490213

```

In [47]:

```

clf2 = RandomForestClassifier(max_depth=18, random_state=0)
clf2.fit(X_train, y_train)

```

Out[47]:

```

RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=18, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100,
                        n_jobs=None, oob_score=False, random_state=0, verbose=0,
                        warm_start=False)

```

In [48]:

```
p = clf2.predict(X_test)
count = 0
for i in range(len(p)):
    if(p[i]==y_test[i]):
        count = count + 1
print(count/len(p))
```

0.9999094008535393

In [49]:

```
t.print_confusion_matrix(p,y_test)
```

Out[49]:

| Predicted | 0     | 1      | All    |
|-----------|-------|--------|--------|
| Actual    |       |        |        |
| 0         | 87203 | 16     | 87219  |
| 1         | 3     | 122493 | 122496 |
| All       | 87206 | 122509 | 209715 |

In [ ]:

## Will the flight arrival late before departure or not?

In [50]:

```
fl3 = flights.drop('ARRIVAL_DELAY_STATUS', axis=1)
```

In [51]:

```
x = fl2.values
```

In [52]:

```
Y = flights.ARRIVAL_DELAY_STATUS.values
```

In [53]:

```
pca = PCA()
principalComponents = pca.fit_transform(x)
```

In [54]:

```
X_train, X_test, y_train, y_test = train_test_split(principalComponents, Y, test_size=
0.2, random_state=0)
```

In [55]:

```

maxi = -1
index = -1
for j in range(20):
    if(j == 0):
        continue
    clf = RandomForestClassifier(max_depth=j, random_state=0)
    clf.fit(X_train, y_train)
    p = clf.predict(X_test)
    count = 0
    for i in range(len(p)):
        if(p[i] == y_test[i]):
            count = count + 1
    accu = count/len(p)
    if(maxi < accu):
        maxi = accu
        index = j
    print(j, count/len(p))

```

```

1 0.9273299477862814
2 0.9747895954032854
3 0.990549078511313
4 0.996833798250006
5 0.9979829768972176
6 0.9987363803256801
7 0.9992513649476671
8 0.9994468683689769
9 0.9996137615335098
10 0.9997329709367475
11 0.9997949598264311
12 0.9998378752115967
13 0.9998474119638557
14 0.9998521803399852
15 0.9998664854683738
16 0.9998903273490213
17 0.9998712538445033
18 0.9998807905967623
19 0.9998855589728918

```

In [56]:

```

clf2 = RandomForestClassifier(max_depth=19, random_state=0)
clf2.fit(X_train, y_train)

```

Out[56]:

```

RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=16, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100,
                        n_jobs=None, oob_score=False, random_state=0, verbose=0,
                        warm_start=False)

```

In [57]:

```
p = clf2.predict(X_test)
count = 0
for i in range(len(p)):
    if(p[i]==y_test[i]):
        count = count + 1
print(count/len(p))
```

0.9998903273490213

In [58]:

```
t.print_confusion_matrix(p,y_test)
```

Out[58]:

| Predicted | 0      | 1     | All    |
|-----------|--------|-------|--------|
| Actual    |        |       |        |
| 0         | 130276 | 0     | 130276 |
| 1         | 23     | 79416 | 79439  |
| All       | 130299 | 79416 | 209715 |

In [ ]:

## Predict Value of Taxi Out on the basis of Taxi In (Linear Regression)

In [ ]:

```
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
import matplotlib.pyplot as plt # To visualize
```

In [322]:

```
Xw = fd.iloc[:, 13].values#.reshape(-1, 1) # values converts it into a numpy array
Yw = fd.iloc[:, 18].values#.reshape(-1, 1)
d = fd.dropna()

Xw = d.iloc[:, 13].values#.reshape(-1, 1) # values converts it into a numpy array
Yw = d.iloc[:, 18].values#.reshape(-1, 1)

# Xw = [3 if np.isnan(i) else i for i in Xw]
# Yw = [3 if np.isnan(i) else i for i in Yw]
Xw = np.array(Xw).reshape(-1,1)
Yw = np.array(Yw).reshape(-1,1)
```

In [323]:


```
X_train, X_test, y_train, y_test = train_test_split(Xw, Yw, test_size=0.2, random_state=0)
```

In [324]:

```
linear_regressor = LinearRegression()  
linear_regressor.fit(X_train, y_train)  
Y_pred = linear_regressor.predict(X_train)
```

In [325]:

```
plt.scatter(X_train, y_train)  
plt.plot(X_train, Y_pred, color='b')  
plt.title('(Linear Regression)')  
plt.xlabel('ELAPSED_TIME')  
plt.ylabel('SCHEDULED_TIME')  
plt.show()  
plt.show()
```



In [ ]:

## Relation of Scheduled as Elapsed Time

In [326]:

```
d = fd.dropna()

Xz = d.iloc[:, 14].values#.reshape(-1, 1) # values converts it into a numpy array
Yz = d.iloc[:, 15].values#.reshape(-1, 1)

# Xw = [3 if np.isnan(i) else i for i in Xw]
# Yw = [3 if np.isnan(i) else i for i in Yw]
Xz = np.array(Xz).reshape(-1,1)
Yz = np.array(Yz).reshape(-1,1)
```

In [327]:

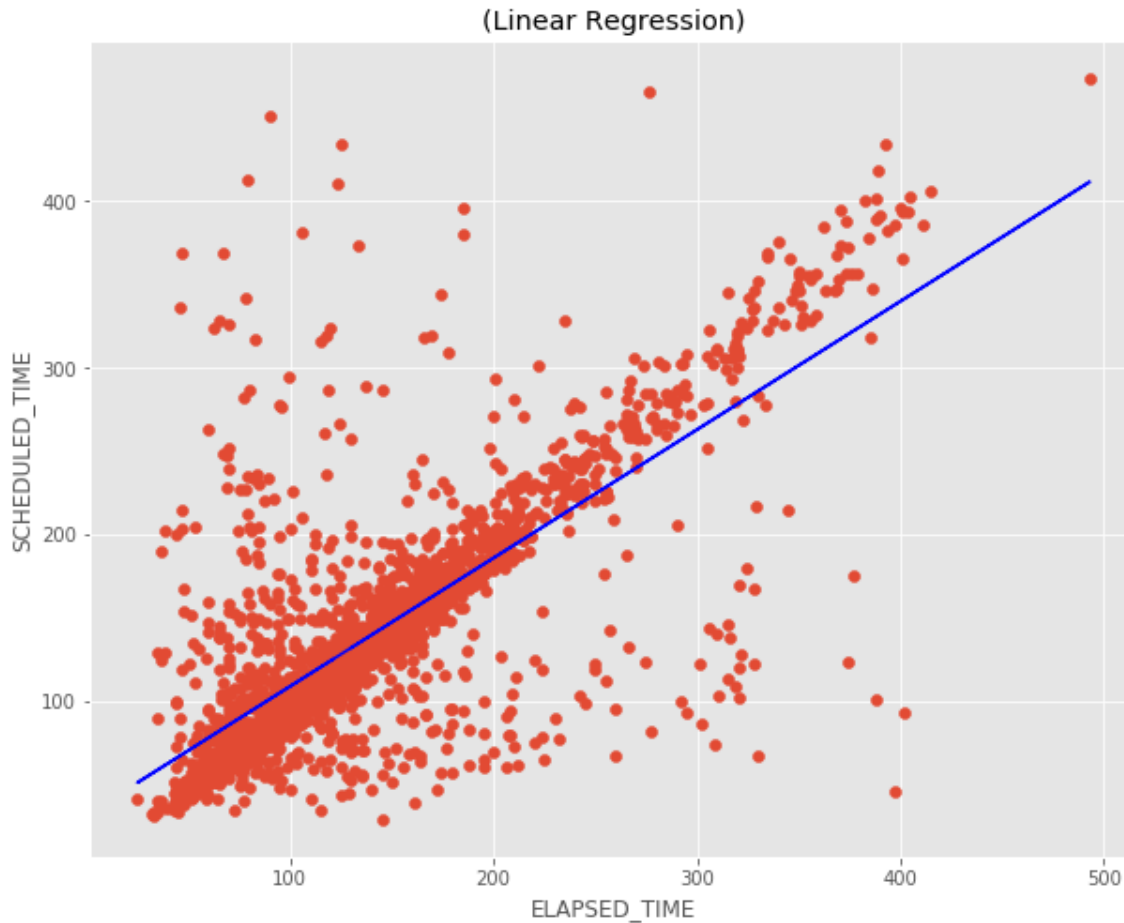
```
X_train, X_test, y_train, y_test = train_test_split(Xz, Yz, test_size=0.2, random_state=0)
```

In [328]:

```
linear_regressor = LinearRegression()
linear_regressor.fit(X_train, y_train)
Y_pred = linear_regressor.predict(X_train)
```

In [329]:

```
plt.scatter(X_trai, y_trai)
plt.plot(X_trai, Y_pred, color='b')
plt.title('(Linear Regression)')
plt.xlabel('ELAPSED_TIME')
plt.ylabel('SCHEDULED_TIME')
plt.show()
```



In [330]:

```
poly_reg=PolynomialFeatures(degree=4)
X_poly=poly_reg.fit_transform(X_trai)
poly_reg.fit(X_poly,y_trai)
lin_reg2=LinearRegression()
lin_reg2.fit(X_poly,y_trai)
```


Out[330]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```



In [331]:

```
plt.scatter(X_train,y_train,color='red')
plt.plot(X_train,lin_reg2.predict(poly_reg.fit_transform(X_train)),color='blue')
plt.title('Polynomial Regression at degree 4')
plt.xlabel('ELAPSED_TIME')
plt.ylabel('SCHEDULED_TIME')
plt.show()
```



In [ ]:

## Check with Different Degrees

In [332]:

```

degrees = [2, 3, 4, 5, 6, 7]

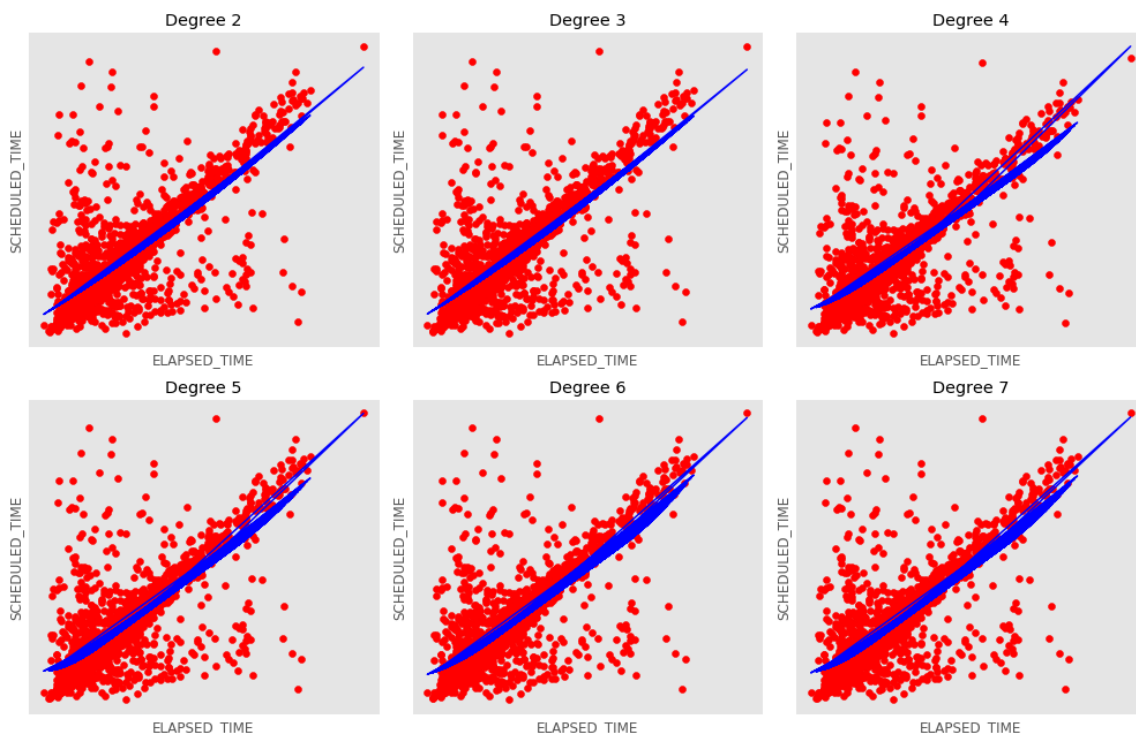
plt.figure(figsize=(14, 9))
for i in range(len(degrees)):
    ax = plt.subplot(2, 3, i + 1)
    plt.setp(ax, xticks=(), yticks=())

    poly_reg=PolynomialFeatures(degree=degrees[i])
    X_poly=poly_reg.fit_transform(X_train)
    poly_reg.fit(X_poly,y_train)
    lin_reg2=LinearRegression()
    lin_reg2.fit(X_poly,y_train)

    plt.scatter(X_train,y_train,color='red')
    plt.plot(X_train,lin_reg2.predict(poly_reg.fit_transform(X_train)),color='blue')
    plt.title("Degree " + str(degrees[i]))
    plt.xlabel('ELAPSED_TIME')
    plt.ylabel('SCHEDULED_TIME')

plt.tight_layout()
plt.show()

```



In [ ]: